

hw5: machine teaching for kNN

March 17, 2023

Similar to hw3, we now consider pool-based teaching for kNN learners.

- Given any dataset $D = (x_1, y_1) \dots (x_n, y_n)$ where $x_i \in \mathbb{R}^d$ and y_i is now a class label, The student runs kNN to learn a classifier $f : X \mapsto Y$. k is given and fixed.
- The teacher knows the above student algorithm. The teacher can only influence the student with the teaching set D . The teacher wants to make sure the student classifier f is close to a target classifier $g : X \mapsto Y$.

1 Formal Definition

Recall a pool-based teacher cannot create arbitrary data points. Instead, the teacher is given a “pool” of data points $P = \{(x_1, y_1) \dots (x_N, y_N)\}$, and the teacher must select pairs from the pool to form its teaching set D . For this homework, we allow D to be a multiset (i.e. allow repeated items from P). Exact teaching is in general infeasible in pool-based teaching. Instead, the teacher aims to approximately teach the student the target model g .

In theory, let

$$f = kNN(D)$$

where we treat kNN as a function (a learning algorithm) that takes a training set D and outputs a classifier f . The teacher has a target classifier g , and has a probability density function $p(x)$ over the input space X . Let the 0-1 loss be $\ell(y, y') = 1$ if $y \neq y'$ and 0 otherwise. Then the teacher can define the disagreement between f and g as

$$d(f, g) := \int p(x) \ell(f(x), g(x)) dx. \quad (1)$$

Clearly, $d(f, g) \in [0, 1]$. $d(f, g) = 0$ if f equals g with probability one. We can now state the teaching problem as

$$D^* = \operatorname{argmin}_{D \subseteq P} d(kNN(D), g). \quad (2)$$

2 Approximating Integral with Sampling

In practice, integrating over X can be difficult. Instead, we approximate integral with sampling. For this, we need a large sample

$$Z = \{(x'_1, y'_1 = g(x'_1)), \dots, (x'_m, y'_m = g(x'_m))\}$$

where $x'_i \sim p(x)$. In general, Z is distinct from the pool P , though the two could be the same. We can then approximate the integral by average and define

$$\hat{d}(f, g) := \frac{1}{m} \sum_{i=1}^m \ell(f(x'_i), y'_i). \quad (3)$$

It will be the case that $\hat{d}(f, g) \approx d(f, g)$ when m is large. Note $\hat{d}(f, g)$ is a random variable because it depends on the sample Z , while $d(f, g)$ is a constant. Given Z , the teacher can solve a slightly different problem

$$\hat{D} = \operatorname{argmin}_{D \subseteq P} \quad \hat{d}(kNN(D), g). \quad (4)$$

Again, there are at least two ways to solve this problem:

1. Enumeration. The teacher enumerates all subsets of P . Note a subset can be even smaller than k : kNN is well defined on any training set size. If a training set is smaller than k , kNN will simply do a majority vote on all training items.
2. Greedy. Given a partial teaching set D (initially empty), the teacher enumerates all one-item extension $D \cup \{(x, y)\}$ where $(x, y) \in P$. The teacher adds the best one-item extension (x^*, y^*) to D , where

$$(x^*, y^*) = \operatorname{argmin}_{(x, y) \in P} \quad \hat{d}(kNN(D \cup \{(x, y)\}), g). \quad (5)$$

This greedy process repeats so D grows.

3 Adding a Teaching Cost

The teacher's objective so far is to guide the student close to g . Finally, we introduce the notation of a "teaching cost" function $c(D)$ which specifies how costly it is for the teacher to use a teaching set D . The new teaching problem is

$$\hat{D} = \operatorname{argmin}_{D \subseteq P} \quad \hat{d}(kNN(D), g) + c(D). \quad (6)$$

$c(D)$ can be thought of as a "regularizer" in the data space that encourages cheap teaching sets. For this homework, we will consider a particularly simple teaching cost: $c(D) = 0$ if $|D| \leq n^*$, and ∞ otherwise, for a given threshold n^* . This simply prevents D from being larger than n^* . For enumeration, you will consider subsets of size at most n^* ; for greedy, simply stop after size n^* .

4 Hand in

For this homework let $k = 1$ (1NN), $n^* = 20$, $P = Z = \text{hw5data.txt}$. Each row $x_1 \ x_2 \ y$ is a data item with 2D feature (x_1, x_2) and binary label y . Implement both enumeration and greedy. For each teaching set size, show:

1. the number of teaching sets of that size that you have to search through;
2. number of seconds it takes for that size;
3. $\hat{d}(kNN(D), g)$ of the best teaching set of that size
4. for enumeration, plot the best teaching set \hat{D} in relation to P (i.e. plot both, but use different symbols for \hat{D})

For greedy, only plot one figure for the last teaching set \hat{D} with size n^* in relation to P , but see if you can mark the order of teaching items entering the greedy teaching set (e.g. use numbers instead of symbols to show items in \hat{D}).

Again you may not be able to enumerate teaching sets of large size, and that is OK.